

Pronto: Trust & Reliability — Product Requirements Document

Type: Feature Proposal (built on top of the existing, live Pronto product) **Status:** Draft v3.1

Author: Product Team **Last updated:** July 25, 2026

1. Problem / Opportunity

Pronto already works at scale: 500,000+ homes served, 10,000+ trained "Pronto Pros," live in 11 Indian cities, with instant/scheduled/recurring bookings and flat, transparent pricing. The core promise — trusted house help, in minutes — is real and differentiated.

But the problem isn't a missing feature. The problem is that **trust breaks down at a few specific, recurring moments in the experience**, and each breakdown undermines the core brand promise in a different way:

- **The clock stops being believable.** When the quoted arrival time doesn't match reality, the "in minutes" promise stops being something customers can plan around — which is the exact pain Pronto exists to remove.
- **The visit is a black box.** When a customer can't see what's happening during a visit, they're back to the "just hope it gets done right" uncertainty they had before Pronto, and they end up policing the work themselves — re-introducing the very effort the product was sold to eliminate.
- **Recovery is harder than the original problem.** When a booking goes wrong, the customer's cost isn't just the failed visit — it's the time spent chasing a resolution. A customer who loses 40 minutes recovering ₹200 has learned that Pronto's failures are expensive to fix, and that lesson outlives the refund.
- **"Pronto" doesn't mean one thing.** When quality and reliability vary sharply by city, the brand becomes a lottery based on where you live, and the national average hides the cities that are failing.
- **Worker safety is structurally unsolved.** Pros — mostly women, spending extended unsupervised time inside strangers' homes — face a real personal-safety risk that today's tooling (an SOS button, self-defense training) only addresses *after* something has already started going wrong. Industry investors and press coverage both point to worker safety as the single biggest determinant of which platform in this space earns durable trust. This isn't a "nice to have" — it's arguably the core long-term risk to the business model itself.

The support insight that reframes this proposal: Most support contacts today are not support problems — they are product gaps in disguise. A customer calls because the app gave them no way to cancel within a grace window, no way to see whether their refund was approved, no way to see how long they'll wait for an agent, and no way to avoid re-explaining the whole story to a second agent. Adding support headcount treats the symptom. The fix is to make the majority of these contacts unnecessary, and to make the remainder resolve on the first touch.

The opportunity: Pronto has already won the "can you get help fast" battle. The next competitive edge — and the thing that will most determine retention, worker supply, and brand durability — is closing the gap between what's promised and what's reliably, verifiably delivered on both sides of the marketplace.

2. Target Users & Use Cases

User	Primary use case this proposal addresses
Repeat household customer (books weekly/recurring help)	Wants confidence that every visit, not just the first one, will be on time, complete, and match their home's known preferences — without re-briefing or re-earning trust each time.
First-time / occasional customer	Deciding whether to trust a stranger in their home for the first time; abandons the platform quickly if the first experience breaks the promise and recovering from it is slow.
Customer whose booking has just failed	In the highest-stakes moment on the platform. Wants the problem fixed or their money back without a phone queue — and this is the moment that decides whether they ever book again.
Pronto Pro (worker)	Wants steady, safe work — specifically, wants to feel protected <i>during</i> a visit, not just have a button to press after something has already gone wrong.
City operations team	Responsible for local quality and Pro supply; needs visibility into where their city is underperforming before it surfaces as public reviews or press coverage — and needs to do reassignments inside the system rather than around it.

3. Current Landscape

What customers experience today: An instant or scheduled booking with a quoted ETA (commonly marketed as ~15 minutes), a fixed price shown up front, and a Pro who arrives, does the job, and leaves — with no live visibility into the visit unless the customer pays for the premium "Pronto Verified" recorded-session add-on (which itself isn't available for recurring bookings). If anything goes wrong, the path is in-app chat, email, or a phone line.

What's happening in practice, based on public app store reviews and press coverage:

Reliability

- Actual arrival times have been reported well beyond the quoted estimate (up to 40–50 minutes against a ~15-minute promise) in some bookings.
- Reports of repeated no-shows and inconsistent training are concentrated in specific (often newer) cities rather than spread evenly.

- Customers report having to direct Pros task by task and finding work claimed as done that wasn't — with no shared record of what was actually booked.

Support and recovery — the sharpest gap

- Support is staffed roughly 7 AM to 10 PM, while booking is marketed as 24/7. A customer with a problem on a late-night or early-morning booking has no live channel at all.
- Reported experience includes long silences mid-conversation (one customer described waiting ~40 minutes with no response), then being connected to a second agent who had no context and had to be re-told the entire issue.
- Agents have been reported as initially saying a cancellation "isn't possible" before negotiating an ad-hoc workaround — suggesting frontline staff either lack the authority or lack a clear policy engine to apply.
- Customers have reported that in-app cancel and reschedule paths show no available slots or a blanket "no refund" message, pushing them into the support queue to claim something the published cancellation policy already entitles them to.
- Refunds are widely described as inconsistently processed after a failed visit, with no visible status once requested.
- Local ops representatives have been reported messaging customers directly and manually reassigning Pros outside the app's auto-assignment, creating confusion and last-minute cancellations — coordination happening around the product rather than inside it.
- Home context and preferences (for example, that there's a pet in the house) don't reliably carry between visits, so customers re-explain their home every time and friction repeats.

Worker safety

- Reuters' April 2026 reporting on India's on-demand house-help sector frames worker safety — given that workers, mostly women, spend hours alone inside customers' homes — as the industry's central unresolved risk, ahead of comparable last-mile-delivery risk, and quotes an investor in a competing platform describing safety as the deciding factor for long-term consumer loyalty.

Competitive landscape: Pronto competes with Snabbit (also on-demand, also cites in-app SOS as a safety measure) and Urban Company (larger, more established, broader service categories, and noted by reviewers for stronger consistency controls such as choosing a preferred professional and tracking service time). None of the three has publicly solved the reliability/recovery/safety gap completely — this is a category-wide problem, which makes solving it well a genuine differentiation opportunity rather than table stakes.

4. Proposed Solution — Elevator Pitch

In one line: Make the promise Pronto already makes — "trusted house help, in minutes" — actually true every time, by closing the gap between what's quoted and what's visible,

verifiable, recoverable, and enforced on both the customer and worker side.

Top 3 value props:

- 1. **Customers can trust the clock** — the ETA they see is the ETA they get, and if it slips, they're compensated automatically rather than after filing a complaint.
- 2. **Customers can see the visit, and fix it themselves when it breaks** — live status and task progress on every booking, plus self-serve cancellation, refund status, and issue reporting that don't require reaching a human at all.
- 3. **Pros are protected proactively, not just reactively** — safety tooling checks in on the worker during a visit, rather than waiting for the worker to signal that something has already gone wrong.

Conceptual model:



5. Goals / Measurable Outcomes

- **Fewer bookings need a human at all:** the share of bookings that generate a support contact falls, because cancellations, reschedules, refund status, and issue reports are all resolvable in-app.
- **Failures stop being expensive to recover from:** median time from "something went wrong" to "resolved or refunded" drops sharply, and no customer waits on an

unanswered channel during an active booking.

- **Trust-related churn falls:** the repeat-booking rate among customers whose first booking failed (late, no-show, or incomplete) rises measurably toward the rate of customers whose first booking went well.
-

6. MVP / Functional Requirements

Bucketed by the user journey each requirement serves. Priorities: **P0** = required for this initiative to be considered shipped; **P1** = important for a genuinely trustworthy experience; **P2** = valuable follow-on. Each requirement is written so it can be verified as done or not done, without prescribing design or implementation.

Journey A: Booking and waiting for a Pro to arrive

- **[P0]** Customer sees an arrival estimate derived from live conditions rather than a fixed marketed number, and is notified whenever that estimate changes materially before arrival.
- **[P0]** If a Pro arrives later than the estimate the customer was shown, compensation is applied automatically — the customer never has to ask for it.
- **[P1]** Before dispatch, a first-time customer or a booking to a new address completes a lightweight identity and address confirmation.
- **[P1]** A customer can cancel or reschedule a booking themselves, at any point the policy allows, and sees the exact amount that will be refunded *before* confirming — with no support contact required.
- **[P1]** Bookings placed within a short window and cancelled almost immediately (before a Pro has been dispatched) are refunded in full, automatically.

Journey B: During the visit

- **[P0]** Every booking — not only the paid "Verified" tier — shows the customer live status: en route, started, in progress, completed, plus progress against the specific tasks booked.
- **[P0]** The Pro sees the full task list and any home-specific notes for the booking before arriving, so the customer does not have to direct the work task by task.
- **[P0]** The Pro receives a discreet, periodic safety check-in during the visit that does not require them to proactively reach for help — distinct from, and in addition to, the existing SOS button.
- **[P1]** For time-based (hourly) bookings, the customer can see elapsed time against the time they paid for.
- **[P1]** Visits flagged as higher-risk (new customer, new address, isolated locality, late hour) are visible to local supervisors so a proactive check can happen without waiting on a distress signal.

Journey C: When something goes wrong — self-serve first

- **[P0]** A customer can report a specific, structured issue (no-show, late arrival, incomplete task, misconduct, unauthorized cash request, damage) directly from the

booking, rather than only through free-text chat.

- [P0] Issues that can be verified against visit data — for example, a no-show past a grace period, or a visit ended with tasks unstarted — are resolved automatically, with the refund or replacement issued without an agent reviewing them.
- [P0] Once a refund is initiated, the customer can see its current status and expected completion at any time, without asking anyone.
- [P0] A Pro with a substantiated report of soliciting unauthorized cash is suspended pending review — a firm, enforced consequence rather than a policy that exists only in the FAQ.
- [P1] Customers have a legitimate, capped in-app way to tip a Pro, so goodwill has a sanctioned channel and any off-app cash request is unambiguously a violation.

Journey D: When something goes wrong — reaching a human

- [P0] A live support channel is reachable during every hour that bookings can be placed. If Pronto sells a 24/7 booking, a customer with a live booking can reach help 24/7.
- [P0] Issues on an *active* booking are routed ahead of post-visit issues — a customer standing in their home waiting for a Pro is not queued behind a billing question from yesterday.
- [P0] When a customer is handed to a second agent, that agent sees the full prior conversation and booking context — the customer never re-explains the issue.
- [P0] A frontline agent can action a refund, a cancellation, or a replacement within published policy without escalating — no customer is told an entitled outcome "isn't possible."
- [P1] While waiting, the customer can see that their request is still open and roughly where it stands, rather than facing silence.
- [P1] Every support interaction ends with a stated outcome and, where relevant, a committed timeline — visible in-app, not only in the chat transcript.

Journey E: Coming back — recurring bookings and repeat quality

- [P0] Home context and customer preferences (for example, pets in the home, access instructions, task-specific expectations) are stored once and carried into every subsequent booking, and are visible to the assigned Pro.
- [P1] Customers on a recurring booking receive a lightweight, ongoing trust signal for each visit — such as timestamped proof of completion — without the cost or consent overhead of a fully recorded "Verified" session every time.
- [P1] Customers can upgrade any single visit within a recurring series to the full "Verified" experience on demand, rather than choosing between recurring and Verified as an all-or-nothing tradeoff.
- [P2] A customer can indicate a preference for a Pro they've had a good experience with, and the system honours it where availability allows.

Journey F: Local operations

- **[P0]** Any reassignment of a Pro to a booking happens inside the system and is reflected to the customer in-app — off-platform coordination (direct messaging, manual swaps) is not a supported path.
- **[P1]** City operations teams see their own on-time rate, no-show rate, and support-contact rate specifically — not just the company-wide average — so a struggling city is caught internally before it shows up in public reviews.
- **[P2]** A new city's demand and marketing push is paced against the size of its certified Pro pool, so early customers in a new city get the same experience as customers in a mature city.

Telemetry (cross-cutting)

- **[P0]** Product team can monitor quoted-versus-actual arrival variance per city, to know whether Journey A is working.
 - **[P0]** Product team can monitor support-contact rate per booking, first-contact resolution rate, and time-to-resolution — split by issue type — to know whether Journeys C and D are working.
 - **[P0]** Product team can monitor substantiated misconduct and cash-demand reports and their resolution time, to know whether enforcement is working.
 - **[P1]** Product team can compare repeat-booking rates between customers whose first booking succeeded and those whose first booking failed, to know whether recovery is actually saving customers.
-

7. Out of Scope

This proposal deliberately does not cover: new service categories, changes to the pricing model or unit economics, expansion into new cities or countries, a B2B/enterprise offering, or replacing the existing Pronto Verified recording product. Where those intersect with this work (for example, pacing a new city launch), the requirement is stated as a constraint rather than a plan.

8. Open Questions & Assumptions

Assumptions this PRD depends on

- Visit telemetry (Pro location, job start/stop, task completion) either exists today or can be captured, since auto-resolution in Journey C depends on it being trustworthy enough to settle a refund without human review.
- The published cancellation policy is what Pronto actually intends to honour, and the gap is enforcement rather than policy.
- Extending live support to match booking hours is an operational cost the business is willing to carry, or booking hours can be narrowed to match support coverage.

Questions to resolve before build

- What proportion of current support contacts are cancellations, refund status checks, and no-shows? This determines how much of the load Journey C actually removes, and therefore how this work should be sized.
 - Who owns the substantiation process for misconduct and cash-demand reports, and what evidence standard triggers a suspension?
 - Why are local ops teams reassigning Pros manually today — is the auto-assignment insufficient, or is it a habit and trust gap? The fix differs sharply depending on the answer.
 - Does auto-compensation for late arrival create a gameable incentive, and what's the abuse threshold?
 - Is timestamped photo proof legally equivalent to recorded video under the same in-home consent requirements that govern Pronto Verified?
-

This PRD is grounded in publicly available information about the live Pronto product (withpronto.com), app store reviews, and press coverage as of July 2026. Internal usage, support ticket, and incident data should validate and refine scope before final prioritization.